

# BITS F464 - Semester 1 - MACHINE LEARNING

---

## ASSIGNMENT 2 – DECISION TREES AND SUPPORT VECTOR MACHINES

*Team number: 13*

---

*Full names of all students in the team: ANIRUDH BAGALKOTKER, KARTIK PANDEY, ADWAIT KULKARNI, JOY SINHA*

---

*Id number of all students in the team: 2021A7PS2682H, 2021A7PS2574H, 2021A7PS2995H, 2021A8PS1606H*

This assignment aims to identify the differences between three Machine Learning models.

## *1. Preprocess and perform exploratory data analysis of the dataset obtained*

### Import Dependencies and Load Dataset

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.graph_objs as go
```

```
In [ ]: # Load the dataset

# Add column names
column_names = [
    "state", "county", "community", "communityname", "fold", "population",
    "racepctblack", "racePctWhite", "racePctAsian", "racePctHispanic", "agePct",
    "agePct12t29", "agePct16t24", "agePct65up", "numUrban", "pctUrban",
    "pctWWage", "pctWFarmSelf", "pctWInvInc", "pctWSocSec", "pctWPubAsst",
    "medFamInc", "perCapInc", "whitePerCap", "blackPerCap", "indianPerCap",
    "OtherPerCap", "HispPerCap", "NumUnderPov", "PctPopUnderPov", "PctLes",
    "PctNotHSGrad", "PctBSorMore", "PctUnemployed", "PctEmploy", "PctEmpl",
    "PctEmplProfServ", "PctOccupManu", "PctOccupMgmtProf", "MalePctDivorc",
    "MalePctNevMarr", "FemalePctDiv", "TotalPctDiv", "PersPerFam", "PctFa",
    "PctKids2Par", "PctYoungKids2Par", "PctTeen2Par", "PctWorkMomYoungKid",
    "NumIlleg", "PctIlleg", "NumImmig", "PctImmigRecent", "PctImmigRec5",
    "PctImmigRec10", "PctRecentImmig", "PctRecImmig5", "PctRecImmig8", "P",
    "PctSpeakEnglOnly", "PctNotSpeakEnglWell", "PctLargHouseFam", "PctLar",
    "PersPerOccupHous", "PersPerOwnOccHous", "PersPerRentOccHous", "PctPe",
    "PctPersDenseHous", "PctHousLess3BR", "MedNumBR", "HousVacant", "PctH",
    "PctHousOwnOcc", "PctVacantBoarded", "PctVacMore6Mos", "MedYrHousBuil
```

```

    "PctHousNoPhone", "PctW0FullPlumb", "OwnOccLowQuart", "OwnOccMedVal",
    "RentLowQ", "RentMedian", "RentHighQ", "MedRent", "MedRentPctHousInc",
    "MedOwnCostPctInc", "MedOwnCostPctIncNoMtg", "NumInShelters", "NumStr",
    "PctForeignBorn", "PctBornSameState", "PctSameHouse85", "PctSameCity8",
    "PctSameState85", "LemasSwornFT", "LemasSwFTPerPop", "LemasSwFTField0",
    "LemasSwFTFieldPerPop", "LemasTotalReq", "LemasTotReqPerPop", "PolicR",
    "PolicPerPop", "RacialMatchCommPol", "PctPolicWhite", "PctPolicBlack",
    "PctPolicHisp", "PctPolicAsian", "PctPolicMinor", "OfficAssgnDrugUnit",
    "NumKindsDrugsSeiz", "PolicAveOTWorked", "LandArea", "PopDens", "PctU",
    "PolicCars", "PolicOperBudg", "LemasPctPolicOnPatr", "LemasGangUnitDe",
    "LemasPctOfficDrugUn", "PolicBudgPerPop", "ViolentCrimesPerPop"
]

# Define missing values
missing_values = ['?']

# Read the dataset
data = pd.read_csv('communities.data', header=None, names=column_names, n
data.head()

```

```
Out[ ]:
```

	state	county	community	communityname	fold	population	householdsize	race
0	8	NaN	NaN	Lakewoodcity	1	0.19	0.33	
1	53	NaN	NaN	Tukwilacity	1	0.00	0.16	
2	24	NaN	NaN	Aberdeentown	1	0.00	0.42	
3	34	5.0	81440.0	Willingborotownship	1	0.04	0.77	
4	42	95.0	6096.0	Bethlehemtownship	1	0.01	0.55	

5 rows × 128 columns

< >

```
In [ ]: # Remove non-predictive columns
df = data.drop(["state", "county", "community", "communityname", "fold"],
df.head()

```

```
Out[ ]:
```

	population	householdsize	racepctblack	racePctWhite	racePctAsian	racePctHisp	ag
0	0.19	0.33	0.02	0.90	0.12	0.17	
1	0.00	0.16	0.12	0.74	0.45	0.07	
2	0.00	0.42	0.49	0.56	0.17	0.04	
3	0.04	0.77	1.00	0.08	0.12	0.10	
4	0.01	0.55	0.02	0.95	0.09	0.05	

5 rows × 123 columns

< >

## Exploratory Data Analysis

```
In [ ]: samples, features = np.shape(df)
df.shape

```

```
Out[ ]: (1994, 123)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1994 entries, 0 to 1993  
Columns: 123 entries, population to ViolentCrimesPerPop  
dtypes: float64(123)  
memory usage: 1.9 MB
```

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	population	householdsize	racepctblack	racePctWhite	racePctAsian	racePct
count	1994.000000	1994.000000	1994.000000	1994.000000	1994.000000	1994.000000
mean	0.057593	0.463395	0.179629	0.753716	0.153681	0.144000
std	0.126906	0.163717	0.253442	0.244039	0.208877	0.230000
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.010000	0.350000	0.020000	0.630000	0.040000	0.010000
50%	0.020000	0.440000	0.060000	0.850000	0.070000	0.040000
75%	0.050000	0.540000	0.230000	0.940000	0.170000	0.160000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 123 columns

<

>

```
In [ ]: df['ViolentCrimesPerPop'].value_counts(normalize=True)
```

```
Out[ ]: ViolentCrimesPerPop  
0.03    0.052156  
0.04    0.046138  
0.06    0.043129  
0.05    0.040120  
0.02    0.037111  
...  
0.79    0.001003  
0.77    0.000502  
0.89    0.000502  
0.94    0.000502  
0.96    0.000502  
Name: proportion, Length: 98, dtype: float64
```

```
In [ ]: df.drop('ViolentCrimesPerPop', axis=1).skew()
```

```
Out[ ]: population      5.063957
householdsize      0.981300
racepctblack       1.863340
racePctWhite      -1.300489
racePctAsian       2.604395
...
PolicOperBudg      4.153317
LemasPctPolicOnPatr -1.589877
LemasGangUnitDeploy 0.221361
LemasPctOfficDrugUn 2.554246
PolicBudgPerPop     3.222583
Length: 122, dtype: float64
```

```
In [ ]: df.corr()
```

```
Out[ ]:
```

	population	householdsize	racepctblack	racePctWhite	racePctAsi
population	1.000000	-0.046148	0.231178	-0.300845	0.181603
householdsize	-0.046148	1.000000	-0.067109	-0.235907	0.201996
racepctblack	0.231178	-0.067109	1.000000	-0.794389	-0.106738
racePctWhite	-0.300845	-0.235907	-0.794389	1.000000	-0.270266
racePctAsian	0.181603	0.201996	-0.106738	-0.270266	1.000000
...	...	...	...	...	...
LemasPctPolicOnPatr	-0.080482	-0.017972	-0.168434	0.125223	0.069012
LemasGangUnitDeploy	0.100012	-0.000784	0.022388	-0.078552	0.139512
LemasPctOfficDrugUn	0.466352	-0.094368	0.260793	-0.276234	0.101812
PolicBudgPerPop	-0.046494	-0.152603	0.045311	-0.014957	-0.024712
ViolentCrimesPerPop	0.367157	-0.034923	0.631264	-0.684770	0.037612

123 rows × 123 columns

<

>

```
In [ ]: fig = go.Figure(go.Heatmap(z=df.corr(), x=df.corr().columns.tolist(), y=df.corr().columns.tolist()))
fig.show()
```

## Data Preprocessing

```
In [ ]: # Check if there are any missing values
missing_values = df.isnull().sum()
columns_with_missing_values = missing_values[missing_values > (samples // 10)]

# Print the number of missing values for each column
for column in columns_with_missing_values:
    print(f"{column}: {missing_values[column]} missing values")

# Drop columns with too many missing values
df = df.drop(columns=columns_with_missing_values)
```

LemasSwornFT: 1675 missing values  
 LemasSwFTPerPop: 1675 missing values  
 LemasSwFTFieldOps: 1675 missing values  
 LemasSwFTFieldPerPop: 1675 missing values  
 LemasTotalReq: 1675 missing values  
 LemasTotReqPerPop: 1675 missing values  
 PolicReqPerOffic: 1675 missing values  
 PolicPerPop: 1675 missing values  
 RacialMatchCommPol: 1675 missing values  
 PctPolicWhite: 1675 missing values  
 PctPolicBlack: 1675 missing values  
 PctPolicHisp: 1675 missing values  
 PctPolicAsian: 1675 missing values  
 PctPolicMinor: 1675 missing values  
 OfficAssgnDrugUnits: 1675 missing values  
 NumKindsDrugsSeiz: 1675 missing values  
 PolicAveOTWorked: 1675 missing values  
 PolicCars: 1675 missing values  
 PolicOperBudg: 1675 missing values  
 LemasPctPolicOnPatr: 1675 missing values  
 LemasGangUnitDeploy: 1675 missing values  
 PolicBudgPerPop: 1675 missing values

```
In [ ]: # Handle missing values (if any) by replacing them with the mean
df.fillna(df.mean(), inplace=True)
df.head()
```

```
Out[ ]:
```

	population	householdsize	racepctblack	racePctWhite	racePctAsian	racePctHisp	ag
0	0.19	0.33	0.02	0.90	0.12	0.17	
1	0.00	0.16	0.12	0.74	0.45	0.07	
2	0.00	0.42	0.49	0.56	0.17	0.04	
3	0.04	0.77	1.00	0.08	0.12	0.10	
4	0.01	0.55	0.02	0.95	0.09	0.05	

5 rows × 101 columns



```
In [ ]: df.describe()
```

Out [ ]:

	population	householdsize	racepctblack	racePctWhite	racePctAsian	racePct
count	1994.000000	1994.000000	1994.000000	1994.000000	1994.000000	1994.000000
mean	0.057593	0.463395	0.179629	0.753716	0.153681	0.144000
std	0.126906	0.163717	0.253442	0.244039	0.208877	0.230000
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.010000	0.350000	0.020000	0.630000	0.040000	0.010000
50%	0.020000	0.440000	0.060000	0.850000	0.070000	0.040000
75%	0.050000	0.540000	0.230000	0.940000	0.170000	0.160000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 101 columns

```

In [ ]: # Performing PCA on the dataset
target_variable = df['ViolentCrimesPerPop']
features = df.drop(['ViolentCrimesPerPop'], axis=1)

# Standardize the features
standardized_features = (features - features.mean()) / features.std()

# Calculate the covariance matrix
cov_matrix = np.cov(standardized_features, rowvar=False)

# Calculate the eigenvectors and eigenvalues
eigenvalues, eigenvectors = np.linalg.eigh(cov_matrix)

# Sort eigenvalues and corresponding eigenvectors in descending order
sorted_indices = np.argsort(eigenvalues)[::-1]
eigenvalues = eigenvalues[sorted_indices]
eigenvectors = eigenvectors[:, sorted_indices]

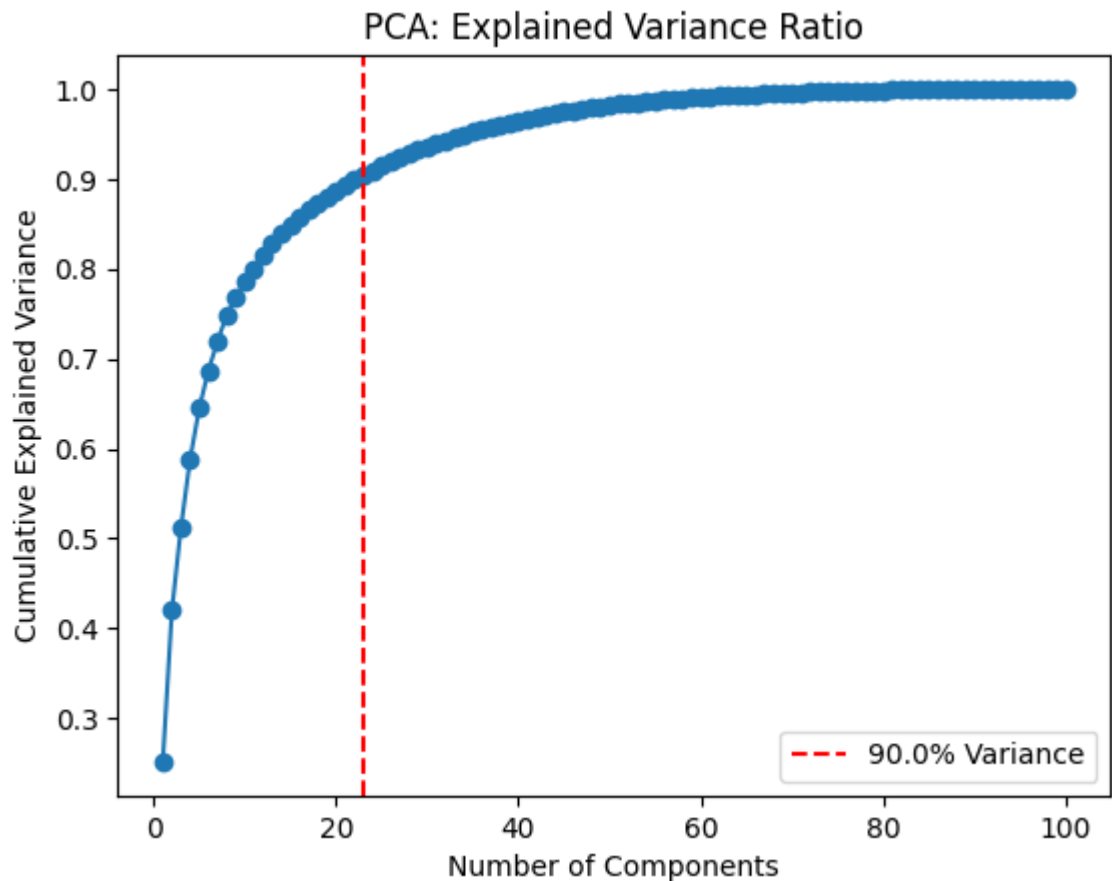
In [ ]: # Calculate the explained variance ratio
explained_variance_ratio = eigenvalues / np.sum(eigenvalues)

# Cumulative explained variance
cumulative_explained_variance = np.cumsum(explained_variance_ratio)

# Find the number of components that explain at least 90% of the variance
desired_explained_variance = 0.90
num_components = np.argmax(cumulative_explained_variance >= desired_explained_variance)

# Plot the explained variance ratio
plt.plot(range(1, len(explained_variance_ratio) + 1), cumulative_explained_variance)
plt.axvline(x=num_components, color='r', linestyle='--', label=f'{desired_explained_variance}')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('PCA: Explained Variance Ratio')
plt.legend()
plt.show()

```



```
In [ ]: # Project the original data onto the selected number of components
selected_eigenvectors = eigenvectors[:, :num_components]
pca_result = np.dot(standardized_features, selected_eigenvectors)

# Create a DataFrame with the principal components and the target variable
df_pca = pd.DataFrame(pca_result, columns=[f'PC{i + 1}' for i in range(num_components)])
df_pca['ViolentCrimesPerPop'] = target_variable
df_pca.head()
df_pca.to_csv('crimes.csv', index=False)
```

## Generate Random Test and Train Splits

```
In [ ]: seed = 420
train_fraction = 0.8
train = df_pca.sample(frac=train_fraction, random_state=seed)
test = df_pca.drop(train.index)
```

```
In [ ]: train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 1595 entries, 493 to 1405
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   PC1                                    1595 non-null   float64
1   PC2                                    1595 non-null   float64
2   PC3                                    1595 non-null   float64
3   PC4                                    1595 non-null   float64
4   PC5                                    1595 non-null   float64
5   PC6                                    1595 non-null   float64
6   PC7                                    1595 non-null   float64
7   PC8                                    1595 non-null   float64
8   PC9                                    1595 non-null   float64
9   PC10                                   1595 non-null   float64
10  PC11                                   1595 non-null   float64
11  PC12                                   1595 non-null   float64
12  PC13                                   1595 non-null   float64
13  PC14                                   1595 non-null   float64
14  PC15                                   1595 non-null   float64
15  PC16                                   1595 non-null   float64
16  PC17                                   1595 non-null   float64
17  PC18                                   1595 non-null   float64
18  PC19                                   1595 non-null   float64
19  PC20                                   1595 non-null   float64
20  PC21                                   1595 non-null   float64
21  PC22                                   1595 non-null   float64
22  PC23                                   1595 non-null   float64
23  ViolentCrimesPerPop  1595 non-null   float64
dtypes: float64(24)
memory usage: 311.5 KB

```

```
In [ ]: test.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Index: 399 entries, 2 to 1989
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PC1                    399 non-null    float64
1   PC2                    399 non-null    float64
2   PC3                    399 non-null    float64
3   PC4                    399 non-null    float64
4   PC5                    399 non-null    float64
5   PC6                    399 non-null    float64
6   PC7                    399 non-null    float64
7   PC8                    399 non-null    float64
8   PC9                    399 non-null    float64
9   PC10                   399 non-null    float64
10  PC11                   399 non-null    float64
11  PC12                   399 non-null    float64
12  PC13                   399 non-null    float64
13  PC14                   399 non-null    float64
14  PC15                   399 non-null    float64
15  PC16                   399 non-null    float64
16  PC17                   399 non-null    float64
17  PC18                   399 non-null    float64
18  PC19                   399 non-null    float64
19  PC20                   399 non-null    float64
20  PC21                   399 non-null    float64
21  PC22                   399 non-null    float64
22  PC23                   399 non-null    float64
23  ViolentCrimesPerPop    399 non-null    float64
dtypes: float64(24)
memory usage: 77.9 KB
```

```
In [ ]: # Assuming 'ViolentCrimesPerPop' is the column you want to predict
X_train = train.drop('ViolentCrimesPerPop', axis=1) # Features for train
y_train = train['ViolentCrimesPerPop'] # Target for training

X_test = test.drop('ViolentCrimesPerPop', axis=1) # Features for testing
y_test = test['ViolentCrimesPerPop'] # Target for testing

# Convert labels to numpy array for applying ML Models
y_train = y_train.to_numpy()
y_test = y_test.to_numpy()
```

```
In [ ]: def accuracy(pred, y_test, threshold=0.5):
    # Calculate the standard deviation of the true values
    y_std = np.std(y_test)

    # Check if the absolute difference is below the threshold (multiple of std)
    correct_predictions = np.abs(pred - y_test) < threshold * y_std

    # Calculate accuracy as the percentage of correct predictions
    accu = 100 * correct_predictions.mean()

    return accu
```

## 2. Decision tree model with entropy implementation

## 2.1 Implementation of the Model

```
In [ ]: class Node():
    def __init__(self, feature_index=None, threshold=None, left=None, right=None, value=None):
        ''' constructor '''

        # for decision node
        self.feature_index = feature_index
        self.threshold = threshold
        self.left = left
        self.right = right
        self.info_gain = info_gain

        # for leaf node
        self.value = value

class DecisionTreeClassifier():
    def __init__(self, min_samples_split=2, max_depth=2):
        ''' constructor '''

        # initialize the root of the tree
        self.root = None

        # stopping conditions
        self.min_samples_split = min_samples_split
        self.max_depth = max_depth

    def build_tree(self, dataset, curr_depth=0):
        ''' recursive function to build the tree '''

        X, Y = dataset[:, :-1], dataset[:, -1]
        num_samples, num_features = np.shape(X)

        # split until stopping conditions are met
        if num_samples >= self.min_samples_split and curr_depth <= self.max_depth:
            # find the best split
            best_split = self.get_best_split(dataset, num_samples, num_features)
            # check if information gain is positive
            if best_split["info_gain"] > 0:
                # recur left
                left_subtree = self.build_tree(best_split["dataset_left"], curr_depth + 1)
                # recur right
                right_subtree = self.build_tree(best_split["dataset_right"], curr_depth + 1)
                # return decision node
                return Node(best_split["feature_index"], best_split["threshold"], left_subtree, right_subtree, best_split["info_gain"])

            # compute leaf node
            leaf_value = self.calculate_leaf_value(Y)
            # return leaf node
            return Node(value=leaf_value)

        # dictionary to store the best split
        best_split = {}
        max_info_gain = -float("inf")
```

```

# loop over all the features
for feature_index in range(num_features):
    feature_values = dataset[:, feature_index]
    possible_thresholds = np.unique(feature_values)
    # loop over all the feature values present in the data
    for threshold in possible_thresholds:
        # get current split
        dataset_left, dataset_right = self.split(dataset, feature
        # check if childs are not null
        if len(dataset_left)>0 and len(dataset_right)>0:
            y, left_y, right_y = dataset[:, -1], dataset_left[:,
            # compute information gain
            curr_info_gain = self.information_gain(y, left_y, rig
            # update the best split if needed
            if curr_info_gain>max_info_gain:
                best_split["feature_index"] = feature_index
                best_split["threshold"] = threshold
                best_split["dataset_left"] = dataset_left
                best_split["dataset_right"] = dataset_right
                best_split["info_gain"] = curr_info_gain
                max_info_gain = curr_info_gain

# return best split
return best_split

def split(self, dataset, feature_index, threshold):
    ''' function to split the data '''

    dataset_left = np.array([row for row in dataset if row[feature_in
    dataset_right = np.array([row for row in dataset if row[feature_i
    return dataset_left, dataset_right

def information_gain(self, parent, l_child, r_child, mode="entropy"):
    ''' function to compute information gain '''

    weight_l = len(l_child) / len(parent)
    weight_r = len(r_child) / len(parent)
    if mode=="gini":
        gain = self.gini_index(parent) - (weight_l*self.gini_index(l_
    else:
        gain = self.entropy(parent) - (weight_l*self.entropy(l_child)
    return gain

def entropy(self, y):
    ''' function to compute entropy '''

    class_labels = np.unique(y)
    entropy = 0
    for cls in class_labels:
        p_cls = len(y[y == cls]) / len(y)
        entropy += -p_cls * np.log2(p_cls)
    return entropy

def gini_index(self, y):
    ''' function to compute gini index '''

    class_labels = np.unique(y)
    gini = 0
    for cls in class_labels:

```

```

        p_cls = len(y[y == cls]) / len(y)
        gini += p_cls**2
    return 1 - gini

def calculate_leaf_value(self, Y):
    ''' function to compute leaf node '''

    Y = list(Y)
    return max(Y, key=Y.count)

def print_tree(self, tree=None, indent=""):
    ''' function to print the tree '''

    if not tree:
        tree = self.root

    if tree.value is not None:
        print(f"{indent}Leaf Node: Class {tree.value}")

    else:
        print(f"{indent}Node: PC{tree.feature_index} <= {tree.threshold}")
        print(f"{indent}left:")
        self.print_tree(tree.left, indent + " ")
        print(f"{indent}right:")
        self.print_tree(tree.right, indent + " ")

def fit(self, X, Y):
    ''' function to train the tree '''

    dataset = np.concatenate((X, Y), axis=1)
    self.root = self.build_tree(dataset)

def predict(self, X):
    ''' function to predict new dataset '''

    predictions = [self.make_prediction(x, self.root) for x in X.values]
    return predictions

def make_prediction(self, x, tree):
    ''' function to predict a single data point '''

    if (tree.value != None):
        return tree.value
    feature_val = x[tree.feature_index]
    if feature_val <= tree.threshold:
        return self.make_prediction(x, tree.left)
    else:
        return self.make_prediction(x, tree.right)

```

```

classifier = DecisionTreeClassifier(min_samples_split=2, max_depth=num_co

```

```

In [ ]: classifier.fit(X_train,pd.DataFrame(y_train))

```

```

In [ ]: Y_pred = np.array(classifier.predict(X_test))
        acc = accuracy(y_test, Y_pred)
        acc

```

```

Out[ ]: 57.89473684210527

```

## ***2.2 Insights drawn (plots, markdown explanations)***

### **Decision Tree Structure**

Lets take a look at the decision tree structure.

```
In [ ]: classifier.print_tree()
```

```

Node: PC0 <= -1.903406546494383 [Info Gain: 0.3068691753511583]
left:
  Node: PC0 <= -5.404072525388851 [Info Gain: 0.2795796614091701]
  left:
    Node: PC10 <= -0.622998399766284 [Info Gain: 0.5069499925013172]
    left:
      Node: PC6 <= 1.304111411122921 [Info Gain: 0.7324452692561572]
      left:
        Node: PC6 <= -0.5355053161787456 [Info Gain: 0.8429038335138657]
        left:
          Node: PC5 <= 0.5314709382419304 [Info Gain: 0.9656361333706105]
          left:
            Node: PC8 <= -0.723970264374518 [Info Gain: 1.0000000000000000]
            left:
              Node: PC0 <= -6.139118842483037 [Info Gain: 0.59167277858232]
              left:
                Node: PC5 <= -0.5463861793856238 [Info Gain: 0.45914791702]
                left:
                  Node: PC3 <= 1.684765404111435 [Info Gain: 0.91829583405]
                  left:
                    Leaf Node: Class 0.66
                  right:
                    Leaf Node: Class 0.36
                right:
                  Leaf Node: Class 0.36
              right:
                Leaf Node: Class 0.93
            right:
              Node: PC0 <= -7.58865331960808 [Info Gain: 0.985228136034251]
              left:
                Node: PC0 <= -8.937900603459456 [Info Gain: 0.918295834054]
                left:
                  Leaf Node: Class 0.61
                right:
                  Node: PC0 <= -7.716220490292774 [Info Gain: 1.0]
                  left:
                    Leaf Node: Class 0.45
                  right:
                    Leaf Node: Class 0.33
              right:
                Node: PC0 <= -6.510808682197184 [Info Gain: 1.0]
                left:
                  Node: PC0 <= -6.8523344043417795 [Info Gain: 1.0]
                  left:
                    Leaf Node: Class 0.72
                  right:
                    Leaf Node: Class 0.54
                right:
                  Node: PC0 <= -5.8161357646481004 [Info Gain: 1.0]
                  left:
                    Leaf Node: Class 0.51
                  right:
                    Leaf Node: Class 0.16
              right:
                right:

```

```

        Leaf Node: Class 0.0
    right:
    Node: PC1 <= 3.9984587146391304 [Info Gain: 1.0]
    left:
    Node: PC1 <= 3.7855241130784365 [Info Gain: 1.0]
    left:
    Leaf Node: Class 0.07
    right:
    Node: PC0 <= 11.788976055157395 [Info Gain: 1.0]
    left:
    Leaf Node: Class 0.05
    right:
    Leaf Node: Class 0.01
    right:
    Node: PC0 <= 8.440018853654108 [Info Gain: 1.0]
    left:
    Node: PC0 <= 7.792827973550439 [Info Gain: 1.0]
    left:
    Leaf Node: Class 0.45
    right:
    Leaf Node: Class 0.03
    right:
    Node: PC0 <= 8.885360375132475 [Info Gain: 1.0]
    left:
    Leaf Node: Class 0.17
    right:
    Leaf Node: Class 0.13

```

## Decision Tree Implementation

### 1. Node Class

The `Node` class represents a node in the decision tree. It has attributes for decision nodes (`feature_index`, `threshold`, `left`, `right`, `info_gain`) and leaf nodes (`value`).

### 2. DecisionTreeClassifier Class

The Decision Tree Classifier is implemented with a recursive binary tree structure. It builds the tree by selecting the best feature and threshold for splitting based on information gain. The tree stops growing when a specified depth or minimum samples for splitting is reached. Leaf nodes represent the majority class, and the structure is printed for interpretability. The classifier is trained using the fit method and makes predictions for new datasets.

#### Initialization:

The class is initialized with parameters `min_samples_split` and `max_depth` to control the tree-building process.

#### Methods:

1. `build_tree` : Recursive tree construction.
2. `get_best_split` : Finds best split based on Entropy.

3. **split** : Divides data based on feature threshold.
4. **information\_gain** : Computes information gain.
5. **calculate\_leaf\_value** : Determines leaf value.
6. **print\_tree** : Prints tree structure.
7. **fit** : Trains the tree.
8. **predict** : Makes predictions.
9. **make\_prediction** : Predicts a single data point.

## Usage

- An instance of DecisionTreeClassifier is created with specified parameters.
- The fit method is called to train the tree on the training data.
- Predictions are made using the predict method on the test data.
- Accuracy is calculated using a simple accuracy calculating function.

## Interpretation of Accuracy

```
In [ ]: print(f"The Accuracy of the Decision Tree Classifier is: {acc}%")
```

The Accuracy of the Decision Tree Classifier is: 57.89473684210527%

## Improvement Suggestions

- Tune hyperparameters.
- Analyze feature importance.
- Use cross-validation.
- Iteratively refine based on insights.
- Advanced Tree algorithms improve implementation and optimization

# 3. Adaboost

## 3.1 Implementation of the Model

```
In [ ]: # Decision stump used as weak classifier
class DecisionStump():
    def __init__(self):
        self.polarity = 1
        self.feature_idx = None
        self.threshold = None
        self.alpha = None

    def predict(self, X):
        n_samples = X.shape[0]
        X_column = X[:, self.feature_idx]
        predictions = np.ones(n_samples)
```



```

    if self.polarity == 1:
        predictions[X_column < self.threshold] = -1
    else:
        predictions[X_column > self.threshold] = -1

    return predictions

```

```

class Adaboost():

```

```

    def __init__(self, n_clf=2):
        self.n_clf = n_clf

    def fit(self, X, y):
        n_samples, n_features = X.shape

        # Initialize weights to 1/N
        w = np.full(n_samples, (1 / n_samples))

        self.clfs = []
        # Iterate through classifiers
        for _ in range(self.n_clf):
            clf = DecisionStump()

            min_error = float('inf')
            # greedy search to find best threshold and feature
            for feature_i in range(n_features):
                X_column = X[:, feature_i]
                thresholds = np.unique(X_column)

                for threshold in thresholds:
                    # predict with polarity 1
                    p = 1
                    predictions = np.ones(n_samples)
                    predictions[X_column < threshold] = -1

                    # Error = sum of weights of misclassified samples
                    misclassified = w[y != predictions]
                    error = sum(misclassified)

                    if error > 0.5:
                        error = 1 - error
                        p = -1

                    # store the best configuration
                    if error < min_error:
                        clf.polarity = p
                        clf.threshold = threshold
                        clf.feature_idx = feature_i
                        min_error = error

            # calculate alpha
            EPS = 1e-10
            clf.alpha = 0.5 * np.log((1.0 - min_error + EPS) / (min_error))

            # calculate predictions and update weights
            predictions = clf.predict(X)

            w *= np.exp(-clf.alpha * y * predictions)
            # Normalize to one

```

```

        w /= np.sum(w)

        # Save classifier
        self.clfs.append(clf)

    def predict(self, X):
        clf_preds = [clf.alpha * clf.predict(X) for clf in self.clfs]
        y_pred = np.sum(clf_preds, axis=0)
        y_pred = np.sign(y_pred)

        return y_pred

classifier = Adaboost()

```

```

In [ ]: classifier.fit(X_train.values, y_train)

y_pred = classifier.predict(X_test.values)

```

```

In [ ]: acc = accuracy(y_test, y_pred, threshold=1)
acc

```

```

Out[ ]: 56.390977443609025

```

## 3.2 Insights drawn (plots, markdown explanations)

### AdaBoost Implementation

#### 1. DecisionStump Class

The `DecisionStump` class represents a weak classifier (a decision stump). It has attributes for polarity, feature index, threshold, and alpha.

#### 2. Adaboost Class

Adaboost, short for Adaptive Boosting, is an ensemble learning algorithm that combines weak classifiers to create a strong classifier. In this implementation, weak classifiers are decision stumps (simple decision trees with a single split). Adaboost iteratively trains decision stumps (`DecisionStump`), adjusting their weights based on their performance. The `Adaboost` Class iteratively selects the best feature and threshold for each weak classifier, assigning higher weights to misclassified samples. The final prediction is a weighted combination of individual weak classifiers. The algorithm adapts by adjusting weights and focuses on difficult-to-classify instances. The resulting ensemble achieves better accuracy than individual classifiers.

##### Initialization:

The class is initialized with the number of weak classifiers (`n_clf`).

##### Methods:

1. `fit`: Trains the AdaBoost ensemble by iteratively training weak classifiers.

2. **predict** : Makes predictions using the ensemble.

## Usage

- An instance of the **AdaBoost** class is created with the specified number of weak classifiers.
- The **fit** method is called to train the AdaBoost ensemble on the training data.
- Predictions are made using the **predict** method on the test data.
- Accuracy is calculated using a simple accuracy calculating function.

## Interpretation of Accuracy

```
In [ ]: print(f"The Accuracy of the AdaBoost Classifier is: {acc}%")
```

The Accuracy of the AdaBoost Classifier is: 56.390977443609025%

## Improvement Suggestions

- Experiment with different weak classifiers.
- Fine-tune hyperparameters, especially the learning rate.
- Consider increasing the number of weak classifiers.
- Evaluate performance on a variety of datasets to ensure generalization.
- Analyze misclassifications for further insights.

# 4. Multiclass SVM

## 4.1 Implementation of the Model

```
In [ ]: class MultiClassSVM:
    def __init__(self, C=1.0, learning_rate=0.01, epochs=500):
        self.C = C # Regularization parameter
        self.learning_rate = learning_rate
        self.epochs = epochs
        self.classifiers = []

    def fit(self, X, y):
        unique_classes = np.unique(y)

        for cls in unique_classes:
            binary_labels = np.where(y == cls, 1, -1)
            classifier = self.train_one_class(X, binary_labels)
            self.classifiers.append((cls, classifier))

    def train_one_class(self, X, y):
        m, n = X.shape
        weights = np.zeros(n)
        bias = 0

        for epoch in range(self.epochs):
            for i in range(m):
                if y[i] * (np.dot(X[i], weights) - bias) >= 1:
```

```

        weights -= self.learning_rate * (2 * self.C * weights
else:
        weights -= self.learning_rate * (2 * self.C * weights
        bias -= self.learning_rate * y[i]

    return (weights, bias)

def predict(self, X):
    predictions = []

    for cls, classifier in self.classifiers:
        weights, bias = classifier
        decision = np.dot(X, weights) - bias
        predictions.append((cls, decision))

    # Choose the class with the highest decision value as the predict
    return max(predictions, key=lambda x: x[1])[0]

# Convert labels to binary for each class
def to_binary_labels(y, target_class):
    return np.where(y == target_class, 1, -1)

# Train the SVM classifier
classifier = MultiClassSVM()
classifier.fit(X_train.values, y_train)

# Make predictions
predictions = [classifier.predict(x) for x in X_test.values]

```

```

In [ ]: acc = accuracy(y_test, predictions)
acc

```

```

Out[ ]: 58.64661654135338

```

## 4.2 Insights drawn (plots, markdown explanations)

### Multi-Class SVM Implementation

#### to\_binary\_labels Function

The `to_binary_labels` function converts the multi-class labels to binary labels for each class, where the target class is assigned a label of 1 and all other classes are assigned a label of -1.

#### MultiClassSVM Class

The `MultiClassSVM` class implements a multi-class Support Vector Machine (SVM) using a **ONE-VS-ALL** strategy. It has attributes for the regularization parameter ( `C` ), learning rate ( `learning_rate` ), and number of epochs ( `epochs` ). The trained classifiers for each class are stored in the `classifiers` attribute.

**Methods:**

1. **fit Method:** Iterates over unique classes, converts labels to binary, and trains a binary SVM for each class. The trained classifiers are stored in the `classifiers` attribute.
2. **train\_one\_class Method:** Trains a binary SVM for one class using stochastic gradient descent (SGD) with hinge loss.
3. **predict Method:** Makes predictions by obtaining decision values for each class and selecting the class with the highest decision value as the prediction.

## Usage

- An instance of the `MultiClassSVM` class is created with specified parameters.
- The `fit` method is called to train the multi-class SVM on the training data.
- Predictions are made using the `predict` method on the test data.
- Accuracy is calculated using a simple accuracy calculating function.

## Interpretation of Accuracy

```
In [ ]: print(f"The Accuracy of the Multi-Class SVM Classifier is: {acc}%")
```

The Accuracy of the Multi-Class SVM Classifier is: 58.64661654135338%

## Improvement Suggestions

- Fine-tune hyperparameters (e.g., `C`, `learning_rate`, `epochs`).
- Evaluate performance on various datasets to ensure generalization.
- Implement kernelized SVM for non-linear decision boundaries.
- Explore additional multi-class SVM strategies (e.g., one-vs-one).
- Use Cross-validation strategy to evaluate the Model using standard SVM libraries.

## 5. References

1. Dataset Description - <https://www.hindawi.com/journals/cin/2022/9283293/>
2. EDA and Data Cleaning - <https://www.kaggle.com/code/charmainechiam/dealing-with-missing-values-in-data-preparation>
3. PCA without sklearn: <https://towardsdatascience.com/principal-component-analysis-pca-from-scratch-in-python-7f3e2a540c51>
4. SVM: <https://machinelearningmastery.com/one-vs-rest-and-one-vs-one-for-multi-class-classification/>
5. AdaBoost: [https://www.python-engineer.com/courses/mlfromscratch/13\\_adaboost/](https://www.python-engineer.com/courses/mlfromscratch/13_adaboost/)
6. Decision Trees: <https://www.analyticsvidhya.com/blog/2020/10/all-about-decision-tree-from-scratch-with-python-implementation/>

